

Introduction:

Our project embarks on training a YOLOv8 model, the pinnacle of real-time object detection. YOLO, short for "You Only Look Once," revolutionized object detection by prioritizing speed and efficiency without compromising accuracy. YOLOv8 continues this legacy, striking a balance between speed and precision.

Understanding YOLO Models:

Innovation in Object Detection: YOLO's core principle is fast, single-shot object detection. It swiftly processes images, making it ideal for applications like autonomous vehicles and surveillance.

YOLOv8: The Evolution: We focus on YOLOv8, the latest YOLO iteration. It merges past advancements, offering a versatile solution for various domains.

Custom Dataset Training: We'll train YOLOv8 on a custom dataset, tailored to our objectives. This involves data collection, annotation, and model training.

Project Goals:

- Train YOLOv8 to detect "Drones" in an image.
- Evaluate the model's performance.
- Fine-tune for optimal results.

In this project, we explore YOLOv8's capabilities and leverage deep learning to create a customized, accurate, and efficient object detection solution.

"YOLOv8, the most recent addition to the YOLO family of real-time object detectors, combines state-of-the-art accuracy and speed. It builds upon the innovations from earlier YOLO versions, introducing novel features and optimizations that render it a top pick for a multitude of object detection tasks across diverse applications."




image source : [Ultralytics](#)

```
# Importing required packages
!pip install ultralytics
```

```
import os
import shutil
from ultralytics import YOLO
import random
import cv2
import matplotlib.pyplot as plt
import numpy as np
```

In the first step I will download a custom dataset from [Roboflow](#) in YOLOv8 format for our task which is "drone" detection.

```
# Downloading the dataset
!curl -L https://universe.roboflow.com/ds/[REDACTED] >
/content/roboflow.zip
```

% Total	% Received	% Xferd	Average	Speed	Time	Time	Time
Current			Dload	Upload	Total	Spent	Left
Speed							
100 894 100 894 0 0 2577 0	--:--:--	--:--:--					
--:--:-- 2576							
100 49.3M 100 49.3M 0 0 49.8M 0	--:--:--	--:--:--					
--:--:-- 49.8M							

```
# Creating a directory for our images dataset
!mkdir images

# Unzipping the downloaded file
!unzip /content/roboflow.zip -d /content/images

# Checking the content of our image directory
os.listdir('/content/images')
```

```
['README.roboflow.txt',
 'train',
 'README.dataset.txt',
 'data.yaml',
 'test',
 'valid']
```

Now I'll show 5 random images from our training set next to labeled images with bounding boxes. The YOLOv8 format typically has bounding box coordinates normalized relative to the width and height of the image, in the format [class, x_center, y_center, width, height]. We need to convert these normalized values to actual pixel coordinates for drawing.

```
# Set the paths to the directories
image_dir = '/content/images/train/images'
bbox_dir = '/content/images/train/labels'

# Function to read and convert bounding box coordinates from a file
```

```

def read_bboxes(file_path, img_shape):
    with open(file_path, 'r') as file:
        bboxes = []
        for line in file:
            # YOLO format: class, x_center, y_center, width, height
            (normalized)
            class_id, x_center, y_center, width, height = map(float,
line.strip().split())
            x_center, y_center, width, height = x_center *
img_shape[1], y_center * img_shape[0], width * img_shape[1], height *
img_shape[0]
            x, y = int(x_center - width / 2), int(y_center - height /
2)
            bboxes.append([int(x), int(y), int(width), int(height)])
    return bboxes

# Function to draw bounding boxes on an image
def draw_bboxes(image, bboxes):
    for bbox in bboxes:
        x, y, w, h = bbox
        cv2.rectangle(image, (x, y), (x + w, y + h), (0, 255, 0), 2)
    return image

# Get a list of image file names
image_files = [f for f in os.listdir(image_dir) if f.endswith('.jpg')]
random.shuffle(image_files)

# Select 5 random images
selected_images = image_files[:5]

# Create subplots
fig, axes = plt.subplots(5, 2, figsize=(10, 20))
fig.suptitle('Original and Labeled Images')

for i, img_file in enumerate(selected_images):
    img_path = os.path.join(image_dir, img_file)
    bbox_path = os.path.join(bbox_dir, os.path.splitext(img_file)[0] +
'.txt')

    # Read image
    image = cv2.imread(img_path)
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Convert from BGR
to RGB
    img_shape = image.shape

    # Read and draw bounding boxes
    bboxes = read_bboxes(bbox_path, img_shape)
    image_with_bboxes = draw_bboxes(image.copy(), bboxes)

    # Show original and labeled images

```

```
axes[i, 0].imshow(image)
axes[i, 0].axis('off')
axes[i, 0].set_title('Original')

axes[i, 1].imshow(image_with_bboxes)
axes[i, 1].axis('off')
axes[i, 1].set_title('Image with bounding box')

plt.tight_layout()
plt.show()
```

Original



Original and Labeled Images

Image with bounding box



Original



Image with bounding box



Original



Image with bounding box




```

0 |
| N/A    35C    P8          9W / 70W |          0MiB / 15360MiB |
0%    Default |
|
N/A |
+-----+
+-----+

+-----+
-----+
| Processes:
|
| GPU    GI    CI          PID    Type    Process name
GPU Memory |
|          ID    ID
Usage      |
|
=====
=====|
| No running processes found
|
+-----+
-----+

```

Here just make sure to pass the path of configured data.yaml file for your project!!!!

```
result = model.train(data='/content/data.yaml', epochs=20)
```

```

Ultralytics YOLOv8.0.228 🚀 Python-3.10.12 torch-2.1.0+cu121 CUDA:0
(Tesla T4, 15102MiB)
engine/trainer: task=detect, mode=train, model=yolov8m.pt,
data=/content/data.yaml, epochs=20, time=None, patience=50, batch=16,
imgsz=640, save=True, save_period=-1, cache=False, device=None,
workers=8, project=None, name=train, exist_ok=False, pretrained=True,
optimizer=auto, verbose=True, seed=0, deterministic=True,
single_cls=False, rect=False, cos_lr=False, close_mosaic=10,
resume=False, amp=True, fraction=1.0, profile=False, freeze=None,
overlap_mask=True, mask_ratio=4, dropout=0.0, val=True, split=val,
save_json=False, save_hybrid=False, conf=None, iou=0.7, max_det=300,
half=False, dnn=False, plots=True, source=None, vid_stride=1,
stream_buffer=False, visualize=False, augment=False,
agnostic_nms=False, classes=None, retina_masks=False, show=False,
save_frames=False, save_txt=False, save_conf=False, save_crop=False,
show_labels=True, show_conf=True, show_boxes=True, line_width=None,
format=torchscript, keras=False, optimize=False, int8=False,
dynamic=False, simplify=False, opset=None, workspace=4, nms=False,
lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005,
warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=7.5,
cls=0.5, dfl=1.5, pose=12.0, kobj=1.0, label_smoothing=0.0, nbs=64,

```

optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...

optimizer: AdamW(lr=0.002, momentum=0.9) with parameter groups 77 weight(decay=0.0), 84 weight(decay=0.0005), 83 bias(decay=0.0) 20 epochs...

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
1/20	7.06G	1.552	1.914	1.845	24
640: 100%	██████████	64/64 [00:38<00:00, 1.68it/s]			
	Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100%	██████████	10/10 [00:06<00:00, 1.62it/s]		
	all	293	323	0.016	0.161
0.0059	0.00145				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
2/20	7.26G	1.811	1.946	2.048	28
640: 100%	██████████	64/64 [00:34<00:00, 1.87it/s]			
	Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100%	██████████	10/10 [00:05<00:00, 1.93it/s]		
	all	293	323	0.0103	0.393
0.00691	0.00171				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
3/20	7.29G	1.801	1.912	2.055	30
640: 100%	██████████	64/64 [00:33<00:00, 1.90it/s]			
	Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100%	██████████	10/10 [00:05<00:00, 1.88it/s]		
	all	293	323	0.0201	0.511
0.0148	0.00432				

Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
4/20	7.28G	1.73	1.858	2.02	26
640: 100%	██████████	64/64 [00:33<00:00, 1.94it/s]			
	Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100%	██████████	10/10 [00:05<00:00, 1.75it/s]		

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
	16/20	7.29G	1.231	0.939	1.662	13
640:	100% ██████████	64/64	[00:33<00:00, 1.88it/s]			
		Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100% ██████████	10/10	[00:04<00:00, 2.09it/s]			
		all	293	323	0.895	0.854
0.919	0.543					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
	17/20	7.28G	1.2	0.8922	1.662	11
640:	100% ██████████	64/64	[00:33<00:00, 1.88it/s]			
		Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100% ██████████	10/10	[00:05<00:00, 1.74it/s]			
		all	293	323	0.883	0.772
0.887	0.514					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
	18/20	7.3G	1.169	0.8488	1.632	11
640:	100% ██████████	64/64	[00:33<00:00, 1.90it/s]			
		Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100% ██████████	10/10	[00:05<00:00, 1.78it/s]			
		all	293	323	0.931	0.873
0.94	0.578					

Size	Epoch	GPU_mem	box_loss	cls_loss	df_l_loss	Instances
	19/20	7.29G	1.135	0.7951	1.6	11
640:	100% ██████████	64/64	[00:33<00:00, 1.90it/s]			
		Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100% ██████████	10/10	[00:04<00:00, 2.13it/s]			

		all	293	323	0.894	0.885
0.94	0.581					

Size	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances
------	-------	---------	----------	----------	----------	-----------

	20/20	7.3G	1.115	0.7629	1.574	11
640: 100%	██████████ 64/64 [00:33<00:00, 1.91it/s]					

	Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100%	██████████	10/10	[00:05<00:00, 1.71it/s]	

		all	293	323	0.901	0.9
0.94	0.588					

20 epochs completed in 0.254 hours.

Optimizer stripped from runs/detect/train/weights/last.pt, 52.0MB

Optimizer stripped from runs/detect/train/weights/best.pt, 52.0MB

Validating runs/detect/train/weights/best.pt...

Ultralytics YOLOv8.0.228 🚀 Python-3.10.12 torch-2.1.0+cu121 CUDA:0 (Tesla T4, 15102MiB)

Model summary (fused): 218 layers, 25840339 parameters, 0 gradients, 78.7 GFLOPs

	Class	Images	Instances	Box(P	R
mAP50	mAP50-95): 100%	██████████	10/10	[00:09<00:00, 1.02it/s]	

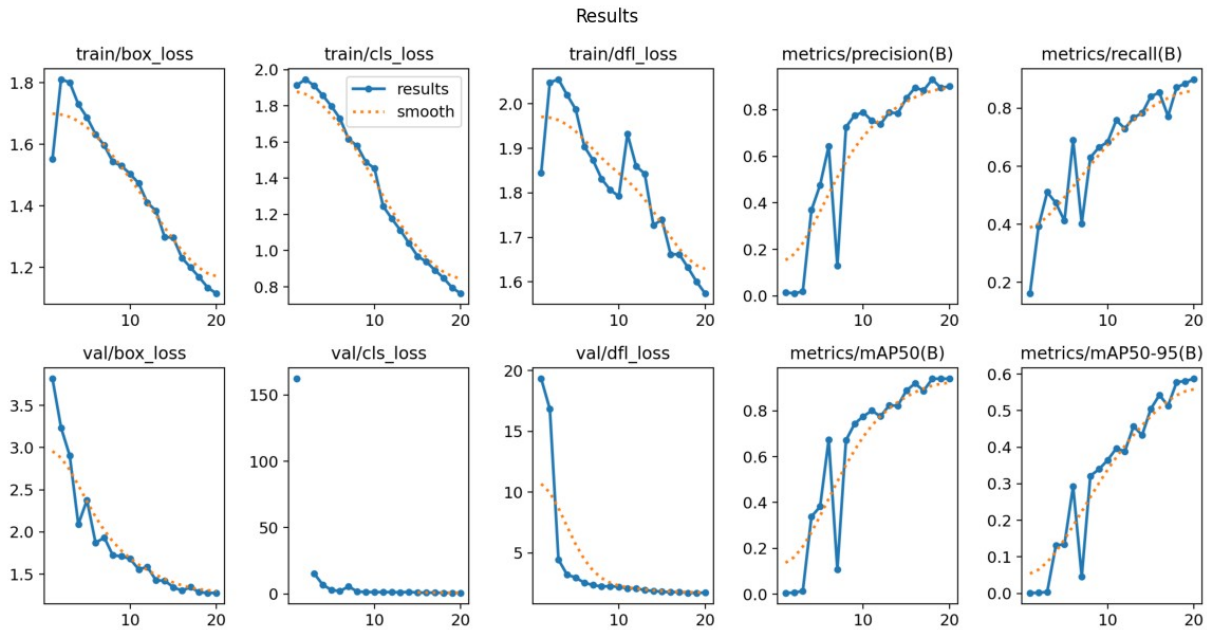
		all	293	323	0.901	0.898
0.94	0.587					

Speed: 0.4ms preprocess, 13.4ms inference, 0.0ms loss, 3.7ms postprocess per image

Results saved to runs/detect/train

```
result_curves =  
cv2.cvtColor(cv2.imread('/content/runs/detect/train/results.png'),cv2.  
COLOR_BGR2RGB)  
plt.figure(figsize=(15,10))  
plt.title('Results')  
plt.axis('off')  
plt.imshow(result_curves)
```

<matplotlib.image.AxesImage at 0x78fd67d1b7c0>



Predicting with Our model

```
results = model("/content/images/test/images", save=True)
```

Now we show some random images from test set with ground truth bounding boxes drawn on it and also the model prediction next to it.

```
# Set the paths to the directories
image_dir = '/content/images/test/images'
bbox_dir = '/content/images/test/labels'

# Function to read and convert bounding box coordinates from a file
def read_bboxes(file_path, img_shape):
    with open(file_path, 'r') as file:
        bboxes = []
        for line in file:
            # YOLO format: class, x_center, y_center, width, height
            # (normalized)
            class_id, x_center, y_center, width, height = map(float,
line.strip().split())
            x_center, y_center, width, height = x_center *
img_shape[1], y_center * img_shape[0], width * img_shape[1], height *
img_shape[0]
            x, y = int(x_center - width / 2), int(y_center - height /
2)
            bboxes.append([int(x), int(y), int(width), int(height)])
    return bboxes

# Function to draw bounding boxes on an image
```

```

def draw_bboxes(image, bboxes):
    for bbox in bboxes:
        x, y, w, h = bbox
        cv2.rectangle(image, (x, y), (x + w, y + h), (0, 255, 0), 2)
    return image

# Get a list of image file names
image_files = [f for f in os.listdir(image_dir) if f.endswith('.jpg')]
random.shuffle(image_files)

# Select 5 random images
selected_images = image_files[:5]

# Create subplots
fig, axes = plt.subplots(5, 2, figsize=(10, 20))
fig.suptitle('Predictions and Ground Truth')

for i, img_file in enumerate(selected_images):
    img_path = os.path.join(image_dir, img_file)
    bbox_path = os.path.join(bbox_dir, os.path.splitext(img_file)[0] +
                              '.txt')

    # Read image
    image = cv2.imread(img_path)
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB) # Convert from BGR
to RGB
    img_shape = image.shape

    # Read and draw bounding boxes
    bboxes = read_bboxes(bbox_path, img_shape)
    image_with_bboxes = draw_bboxes(image.copy(), bboxes)
    prediction =
cv2.cvtColor(cv2.imread('/content/runs/detect/preds/'+img_file),
cv2.COLOR_BGR2RGB)
    # Show original and labeled images
    axes[i, 0].imshow(prediction)
    axes[i, 0].axis('off')
    axes[i, 0].set_title('prediction')

    axes[i, 1].imshow(image_with_bboxes)
    axes[i, 1].axis('off')
    axes[i, 1].set_title('Ground Truth label')

plt.tight_layout()
plt.show()

```


prediction

Predictions and Ground Truth

Ground Truth label

INNLOI UAV

INNLOI UAV



prediction

Ground Truth label



prediction

Ground Truth label

